

## VGG19 With First 5 Layers Frozen

Reference:<https://www.kaggle.com/atrissaxena/keras-plant-seedlings-vgg19-augmentation>

In [1]:

```
import warnings
warnings.filterwarnings('ignore')
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
print(os.listdir("../input"))
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

import cv2
from keras.preprocessing.image import ImageDataGenerator
from PIL import Image
import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid
import numpy as np
from keras.utils import np_utils
from keras import applications
from keras.preprocessing.image import ImageDataGenerator
from keras import optimizers
from keras.models import Sequential, Model
from keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from keras.callbacks import ModelCheckpoint, LearningRateScheduler, TensorBoard, EarlyStopping
```

```
['plant-seedlings-classification', 'withvalid']
```

```
Using TensorFlow backend.
```

In [2]:

```
CATEGORIES = ['Black-grass', 'Charlock', 'Cleavers', 'Common Chickweed', 'Common wheat', 'Fat Hen', 'Loose Silky-bent', 'Maize', 'Scentless Mayweed', 'Shepherds Purse', 'Small-flowered Cranesbill', 'Sugar beet']
```

```
NUM_CATEGORIES = len(CATEGORIES)
```

In [3]:

```
SEED = 123
data_dir = '../input/withvalid'
train_dir = os.path.join(data_dir, 'train_valid/train_valid/train')
valid_dir = os.path.join(data_dir, 'train_valid/train_valid/valid/valid')
test_dir = os.path.join(data_dir, 'test/test')
```

### Number of training images for each Category

In [4]:

```
for category in CATEGORIES:
    print('{} {} images'.format(category, len(os.listdir(os.path.join(train_dir, category)))))

Black-grass 237 images
Charlock 351 images
Cleavers 258 images
Common Chickweed 550 images
Common wheat 199 images
Fat Hen 428 images
Loose Silky-bent 589 images
Maize 199 images
Scentless Mayweed 465 images
Shepherds Purse 208 images
Small-flowered Cranesbill 447 images
Sugar beet 347 images
```

```
In [5]:  
for category in CATEGORIES:  
    print('{} {} images'.format(category, len(os.listdir(os.path.join(valid_dir, category)))))
```

```
Black-grass 26 images  
Charlock 39 images  
Cleavers 29 images  
Common Chickweed 61 images  
Common wheat 22 images  
Fat Hen 47 images  
Loose Silky-bent 65 images  
Maize 22 images  
Scentless Mayweed 51 images  
Shepherds Purse 23 images  
Small-flowered Cranesbill 49 images  
Sugar beet 38 images
```

```
In [6]:  
train = []  
for category_id, category in enumerate(CATEGORIES):  
    for file in os.listdir(os.path.join(train_dir, category)):  
        train.append(['train/{}/{}'.format(category, file), category_id, category])  
train = pd.DataFrame(train, columns=['file', 'category_id', 'category'])  
train.head(2)  
train.shape
```

Out[6]:

|   | file                            | category_id | category    |
|---|---------------------------------|-------------|-------------|
| 0 | train/Black-grass/d3c72d4c3.png | 0           | Black-grass |
| 1 | train/Black-grass/37d85d833.png | 0           | Black-grass |

```
Out[6]:  
(4278, 3)
```

```
In [7]:  
valid = []  
for cat_id, cat in enumerate(CATEGORIES):  
    for f in os.listdir(os.path.join(valid_dir, cat)):  
        valid.append(['valid/{}/{}'.format(cat, file), cat_id, cat])  
valid = pd.DataFrame(valid, columns=['file', 'category_id', 'category'])  
valid.head(5)
```

Out[7]:

|   | file                            | category_id | category    |
|---|---------------------------------|-------------|-------------|
| 0 | valid/Black-grass/dc39e01ec.png | 0           | Black-grass |
| 1 | valid/Black-grass/dc39e01ec.png | 0           | Black-grass |
| 2 | valid/Black-grass/dc39e01ec.png | 0           | Black-grass |
| 3 | valid/Black-grass/dc39e01ec.png | 0           | Black-grass |
| 4 | valid/Black-grass/dc39e01ec.png | 0           | Black-grass |

```
In [8]:  
test = []  
for file in os.listdir(test_dir):  
    test.append(['test/{}'.format(file), file])  
test = pd.DataFrame(test, columns=['filepath', 'file'])  
test.head(2)  
test.shape
```

Out[8]:

|   | filepath           | file          |
|---|--------------------|---------------|
| 0 | test/0e8492cb1.png | 0e8492cb1.png |
| 1 | test/b687160f5.png | b687160f5.png |

Out[8]:

(794, 2)

### See some of the Images

In [9]:

```
fig = plt.figure(1, figsize=(NUM_CATEGORIES, NUM_CATEGORIES))
grid = ImageGrid(fig, 111, nrows_ncols=(NUM_CATEGORIES, NUM_CATEGORIES), axes_pad=0.05)
i = 0
for category_id, category in enumerate(CATEGORIES):
    for filepath in train[train['category'] == category]['file'].values[:NUM_CATEGORIES]:
        ax = grid[i]
        img = Image.open("../input/withvalid/train_valid/train_valid/"+filepath)
        img = img.resize((240,240))
        ax.imshow(img)
        ax.axis('off')
        if i % NUM_CATEGORIES == NUM_CATEGORIES - 1:
            ax.text(250, 112, filepath.split('/')[1], verticalalignment='center')
        i += 1
plt.show();
```



```
In [10]:
model = applications.VGG19(weights = "imagenet", include_top=False, input_shape = (240, 240, 3))
)
for layer in model.layers[:5]:
    layer.trainable = False

x = model.output
x = Flatten()(x)
x = Dense(1024, activation="relu")(x)
x = Dropout(0.5)(x)
x = Dense(1024, activation="relu")(x)
predictions = Dense(12, activation="softmax")(x)

model_final = Model(input = model.input, output = predictions)

model_final.compile(loss = "categorical_crossentropy", optimizer = optimizers.SGD(lr=0.0001, momentum=0.9), metrics=["accuracy"])
model_final.summary()


```

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg19\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5  
80142336/80134624 [=====] - 1s 0us/step

| Layer (type)               | Output Shape          | Param #  |
|----------------------------|-----------------------|----------|
| input_1 (InputLayer)       | (None, 240, 240, 3)   | 0        |
| block1_conv1 (Conv2D)      | (None, 240, 240, 64)  | 1792     |
| block1_conv2 (Conv2D)      | (None, 240, 240, 64)  | 36928    |
| block1_pool (MaxPooling2D) | (None, 120, 120, 64)  | 0        |
| block2_conv1 (Conv2D)      | (None, 120, 120, 128) | 72956    |
| block4_conv1 (Conv2D)      | (None, 30, 30, 512)   | 1180160  |
| block4_conv2 (Conv2D)      | (None, 30, 30, 512)   | 2359808  |
| block4_conv3 (Conv2D)      | (None, 30, 30, 512)   | 2359808  |
| block4_conv4 (Conv2D)      | (None, 30, 30, 512)   | 2359808  |
| block4_pool (MaxPooling2D) | (None, 15, 15, 512)   | 0        |
| block5_conv1 (Conv2D)      | (None, 15, 15, 512)   | 2359808  |
| block5_conv2 (Conv2D)      | (None, 15, 15, 512)   | 2359808  |
| block5_conv3 (Conv2D)      | (None, 15, 15, 512)   | 2359808  |
| block5_conv4 (Conv2D)      | (None, 15, 15, 512)   | 2359808  |
| block5_pool (MaxPooling2D) | (None, 7, 7, 512)     | 0        |
| flatten_1 (Flatten)        | (None, 25088)         | 0        |
| dense_1 (Dense)            | (None, 1024)          | 25691136 |
| dropout_1 (Dropout)        | (None, 1024)          | 0        |
| dense_2 (Dense)            | (None, 1024)          | 1049600  |
| dense_3 (Dense)            | (None, 12)            | 12300    |

Total params: 46,777,420  
Trainable params: 46,664,844  
Non-trainable params: 112,576

```
In [11]:  
    gen = ImageDataGenerator(  
        rotation_range=360.,  
        width_shift_range=0.3,  
        height_shift_range=0.3,  
        zoom_range=0.3,  
        horizontal_flip=True,  
        vertical_flip=True)
```

```
In [12]:  
    train_data_dir = "../input/withvalid/train_valid/train_valid/train"  
    train_generator = gen.flow_from_directory(  
        train_data_dir,  
        target_size=(240, 240),  
        batch_size=16,  
        class_mode="categorical")
```

```
Found 4278 images belonging to 12 classes.
```

```
In [13]:  
    valid_data_dir = "../input/withvalid/train_valid/train_valid/valid/valid"  
    valid_generator = gen.flow_from_directory(  
        valid_data_dir,  
        target_size=(240, 240),  
        batch_size=16,  
        class_mode="categorical")
```

```
Found 472 images belonging to 12 classes.
```

```
In [14]:  
    checkpoint = ModelCheckpoint("vgg16_withvalid.h5", monitor='loss', verbose=1, save_best_only=True,  
    save_weights_only=False, mode='auto', period=1)  
    early = EarlyStopping(monitor='loss', min_delta=0, patience=10, verbose=1, mode='auto')
```

## Train our Model

```
In [15]:  
    batchsize = 16  
    STEP_SIZE_TRAIN=train_generator.n/batchsize  
    STEP_SIZE_VALID=valid_generator.n/batchsize  
  
    model_final.fit_generator(  
        generator=train_generator,  
        validation_data=valid_generator,  
        steps_per_epoch=STEP_SIZE_TRAIN,  
        validation_steps=STEP_SIZE_VALID,  
        epochs=50,  
        shuffle=True,  
        callbacks=[checkpoint, early])
```

```
Epoch 1/50  
268/267 [=====] - 183s 682ms/step - loss: 2.5749 - acc: 0.1203 - val_loss: 2.4212 - val_acc: 0.1547  
  
Epoch 00001: loss improved from inf to 2.57495, saving model to vgg16_withvalid.h5  
Epoch 2/50  
268/267 [=====] - 160s 597ms/step - loss: 2.4085 - acc: 0.1374 - val_loss: 2.3803 - val_acc: 0.1398  
  
Epoch 00002: loss improved from 2.57495 to 2.40887, saving model to vgg16_withvalid.h5  
Epoch 3/50
```

```
Epoch 3/50
268/267 [=====] - 160s 599ms/step - loss: 2.3778 - acc: 0.1574 - val_
loss: 2.3057 - val_acc: 0.1737

Epoch 00003: loss improved from 2.40887 to 2.37705, saving model to vgg16_withvalid.h5
Epoch 4/50
268/267 [=====] - 160s 596ms/step - loss: 2.3057 - acc: 0.1796 - val_
loss: 2.3438 - val_acc: 0.1589

Epoch 00004: loss improved from 2.37705 to 2.30502, saving model to vgg16_withvalid.h5
Epoch 5/50
268/267 [=====] - 159s 592ms/step - loss: 2.0990 - acc: 0.2636 - val_
loss: 1.7064 - val_acc: 0.3983

Epoch 00005: loss improved from 2.30502 to 2.09897, saving model to vgg16_withvalid.h5
Epoch 6/50
268/267 [=====] - 160s 595ms/step - loss: 1.7016 - acc: 0.3932 - val_
loss: 1.3945 - val_acc: 0.4619

Epoch 00006: loss improved from 2.09897 to 1.70268, saving model to vgg16_withvalid.h5
Epoch 7/50
268/267 [=====] - 158s 590ms/step - loss: 1.3843 - acc: 0.4969 - val_
loss: 1.0433 - val_acc: 0.6568

Epoch 00007: loss improved from 1.70268 to 1.38378, saving model to vgg16_withvalid.h5
Epoch 8/50
268/267 [=====] - 158s 591ms/step - loss: 1.0669 - acc: 0.6164 - val_
loss: 0.7816 - val_acc: 0.7564

Epoch 00008: loss improved from 1.38378 to 1.06813, saving model to vgg16_withvalid.h5
Epoch 9/50
268/267 [=====] - 159s 594ms/step - loss: 0.8954 - acc: 0.6847 - val_
loss: 0.8958 - val_acc: 0.6992
```

```
Epoch 00009: loss improved from 1.06813 to 0.89684, saving model to vgg16_withvalid.h5
Epoch 10/50
268/267 [=====] - 160s 598ms/step - loss: 0.7021 - acc: 0.7487 - val_
loss: 0.5431 - val_acc: 0.8114

Epoch 00010: loss improved from 0.89684 to 0.70257, saving model to vgg16_withvalid.h5
Epoch 11/50
268/267 [=====] - 161s 602ms/step - loss: 0.5842 - acc: 0.7937 - val_
loss: 0.4194 - val_acc: 0.8665

Epoch 00011: loss improved from 0.70257 to 0.58469, saving model to vgg16_withvalid.h5
Epoch 12/50
268/267 [=====] - 160s 597ms/step - loss: 0.4928 - acc: 0.8235 - val_
loss: 0.4025 - val_acc: 0.8644

Epoch 00012: loss improved from 0.58469 to 0.49374, saving model to vgg16_withvalid.h5
Epoch 13/50
268/267 [=====] - 160s 598ms/step - loss: 0.4506 - acc: 0.8389 - val_
loss: 0.3796 - val_acc: 0.8750

Epoch 00013: loss improved from 0.49374 to 0.45061, saving model to vgg16_withvalid.h5
Epoch 14/50
268/267 [=====] - 160s 596ms/step - loss: 0.4100 - acc: 0.8553 - val_
loss: 0.4577 - val_acc: 0.8581

Epoch 00014: loss improved from 0.45061 to 0.41028, saving model to vgg16_withvalid.h5
Epoch 15/50
268/267 [=====] - 160s 597ms/step - loss: 0.3682 - acc: 0.8671 - val_
loss: 0.2887 - val_acc: 0.8983

Epoch 00015: loss improved from 0.41028 to 0.36859, saving model to vgg16_withvalid.h5
Epoch 16/50
268/267 [=====] - 160s 597ms/step - loss: 0.3434 - acc: 0.8832 - val_
loss: 0.2457 - val_acc: 0.8750
```

```
#####
Epoch 00038: loss did not improve from 0.14950
Epoch 39/50
268/267 [=====] - 159s 594ms/step - loss: 0.1532 - acc: 0.9459 - val_
loss: 0.2726 - val_acc: 0.9089

Epoch 00039: loss did not improve from 0.14950
Epoch 40/50
268/267 [=====] - 160s 597ms/step - loss: 0.1425 - acc: 0.9459 - val_
loss: 0.1947 - val_acc: 0.9322

Epoch 00040: loss improved from 0.14950 to 0.14275, saving model to vgg16_withvalid.h5
Epoch 41/50
268/267 [=====] - 167s 622ms/step - loss: 0.1340 - acc: 0.9525 - val_
loss: 0.1875 - val_acc: 0.9470

Epoch 00041: loss improved from 0.14275 to 0.13313, saving model to vgg16_withvalid.h5
Epoch 42/50
268/267 [=====] - 160s 596ms/step - loss: 0.1549 - acc: 0.9448 - val_
loss: 0.2178 - val_acc: 0.9386

Epoch 00042: loss did not improve from 0.13313
Epoch 43/50
268/267 [=====] - 160s 597ms/step - loss: 0.1392 - acc: 0.9499 - val_
loss: 0.2089 - val_acc: 0.9301

Epoch 00043: loss did not improve from 0.13313
Epoch 44/50
268/267 [=====] - 159s 594ms/step - loss: 0.1508 - acc: 0.9473 - val_
loss: 0.1819 - val_acc: 0.9470

Epoch 00044: loss did not improve from 0.13313
Epoch 45/50
50/267 [====>.....] - ETA: 1:53 - loss: 0.1530 - acc: 0.9450
```

```
In [18]: prediction = []
for f in test['file']:
    img = cv2.imread(os.path.join(test_dir,f))
    img = cv2.resize(img,(240,240))
    img = np.asarray(img)
    img = img.reshape(1,240,240,3)
    pred = model_final.predict(img)
    prediction.append(classes.get(pred.argmax(axis=-1)[0])) #Invert Mapping helps to map Label
```

```
In [19]: pred = pd.DataFrame({'species': prediction})
test = test.join(pred)
```

### Final submission File

```
In [20]: test.to_csv('submission.csv', index=False)
```

```
In [21]: test.head()
```

Out[21]:

|   | filepath           | file          | species                   |
|---|--------------------|---------------|---------------------------|
| 0 | test/0e8492cb1.png | 0e8492cb1.png | Sugar beet                |
| 1 | test/b687160f5.png | b687160f5.png | Small-flowered Cranesbill |
| 2 | test/fea3da57c.png | fea3da57c.png | Sugar beet                |
| 3 | test/f6d250856.png | f6d250856.png | Shepherds Purse           |
| 4 | test/e73e308be.png | e73e308be.png | Fat Hen                   |

## VGG19+Xception

**The only difference between the combination of pre-trained models and VGG19 with first 5 layers frozen is the model structure as following:**

```
In [8]:  
from keras.layers.convolutional import Conv2D  
from keras.layers.convolutional import MaxPooling2D  
from keras.layers import BatchNormalization  
scale=299  
  
model = applications.VGG19(weights = "imagenet", include_top=False, input_shape = (scale, scale, 3))  
add_model=applications.Xception(weights = "imagenet", include_top=False, input_shape = (scale, scale, 3))  
  
model = Sequential()  
model.add(add_model)  
model.add(Conv2D(20, kernel_size=(3, 3), activation='relu',input_shape=(scale, scale, 3)))  
model.add(BatchNormalization(axis=3))  
model.add(Conv2D(20, kernel_size=(3, 3), activation='relu'))  
model.add(BatchNormalization(axis=3))  
model.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2)))  
model.add(Dropout(0.2))  
  
model.add(Flatten())  
model.add(Dense(256, activation='relu'))  
model.add(Dropout(0.5))  
model.add(Dense(64, activation = 'relu'))  
model.add(Dropout(0.5))  
model.add(Dense(12, activation='softmax'))  
  
#compling and show model  
model.compile(loss = "categorical_crossentropy", optimizer = optimizers.SGD(lr=0.0001, momentum =0.9), metrics=["accuracy"])  
model.summary()
```

```
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/
vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80142336/80134624 [=====] - 1s 0us/step
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.4/
xception_weights_tf_dim_ordering_tf_kernels_notop.h5
83689472/83683744 [=====] - 5s 0us/step
-----
Layer (type)          Output Shape         Param #
=====
xception (Model)     (None, 10, 10, 2048)  20861480
conv2d_5 (Conv2D)    (None, 8, 8, 20)      368660
batch_normalization_5 (Batch (None, 8, 8, 20)  80
conv2d_6 (Conv2D)    (None, 6, 6, 20)      3620
batch_normalization_6 (Batch (None, 6, 6, 20)  80
max_pooling2d_1 (MaxPooling2 (None, 2, 2, 20)  0
dropout_1 (Dropout)  (None, 2, 2, 20)      0
flatten_1 (Flatten)  (None, 80)            0
dense_1 (Dense)     (None, 256)           20736
dropout_2 (Dropout)  (None, 256)           0
dense_2 (Dense)     (None, 64)            16448
dropout_3 (Dropout)  (None, 64)            0
dense_3 (Dense)     (None, 12)             780
dense_3 (Dense)     (None, 12)             780
=====
Total params: 21,271,884
Trainable params: 21,217,276
Non-trainable params: 54,608
```